

Domain	Benchmark	Primary metric	Purpose
Math, moderate	Minerva Math 500 (Lewkowycz et al., 2022)	exact-answer score	Core math reasoning
Math, hard	AIME 2025 (HF set)	pass@1 and pass@32	Hard problems where lengths grow most
Code	LiveCodeBench generation (Jain et al., 2024)	pass@1	Code generation and cross-domain brevity
Instruction following	IFEval / ifbench wildchat + OOD (Zhou et al., 2023)	instruction-following accuracy	Precise instruction following
Chat quality	AlpacaEval v3 (Li et al., 2023)	length-controlled win rate	General quality and verbosity